

OpenStack Telemetry

Rethinking Ceilometer metric storage with Gnocchi: Time-series as a Service

Eoghan Glynn
eglynn@redhat.com
eglynn@Freenode

Julien Danjou
julien@danjou.info
jd__@Freenode

Dina Belova
dbelova@mirantis.com
DinaBelova@Freenode

Speakers



Eoghan Glynn - current Ceilometer PTL, is defining overall vector of project development



Julien Danjou - core Ceilometer contributor from the outset, is driving Gnocchi initiative



Dina Belova - Ceilometer contributor, joined core team during Juno cycle

Once upon a time, Ceilometer...

Let's meter OpenStack!

What do we store?

How do we do store that?

*Let's store everything in some DB,
we'll figure out later!*



How Ceilometer could have been better

- Flexible and heavy Ceilometer samples model with ***free-form metadata***
- Scalability API issue (***queries*** ~ ***O(n)***, where n = number of samples)
- No API consistency
- API cannot solve all problems efficiently

Gnocchi to the rescue



- Let's track **resources**
- Let's ***separate events and metrics*** concepts!
 - Metric = ***Time-Series Data***
- Let's link **resources** and **metrics**
- Let's do **aggregation** eagerly

Compare and contrast ...

“classic” Ceilometer	Gnocchi
Heavy-weight samples with embedded metadata	Light-weight time-series shorn of metadata
Global data expiry policy set across the board	Per time-series configurable retention policies



Compare and contrast ...

“classic” Ceilometer	Gnocchi
On-demand aggregation	Eager pre-aggregation
Intertwined storage of resources and samples	Separated storage and data models for resources & time-series data




Gnocchi basics

- **Resource** = cloud resource (instance, volume, etc.)
- **Entity** = anything you'd like to collect data about
 - identified by UUID, or by name combined with resource ID
- **Measure** = (*timestamp, value*) time-series datapoint

Gnocchi basics

- **Archive policy** = data storage policy defined by admin
 - 1 second resolution over a day, 1 hour resolution over a year, or even both
 - Consists of **granularity** (in seconds) and **retention time-span**
- **Aggregation** = function used to roll up data
- **Retention** = do not store fine grained data forever, instead store aggregated data according to the *per-entity* archive policies

Gnocchi aggregation mechanism



Gnocchi Indexer concept


- Measuring different entities is the main concept of Gnocchi
- Although, entities have no actual use without entity-resource association

Gnocchi Indexer concept

- Gnocchi indexer is responsible for indexing entities, resources, and linking them together
- Resources and their attributes are well-defined, typed, and indexed
 - The generic type can be used if the resource type is unknown to Gnocchi


Gnocchi performance

● ceilometer-collector writes to opentsdb socket via dispatcher




Gnocchi performance

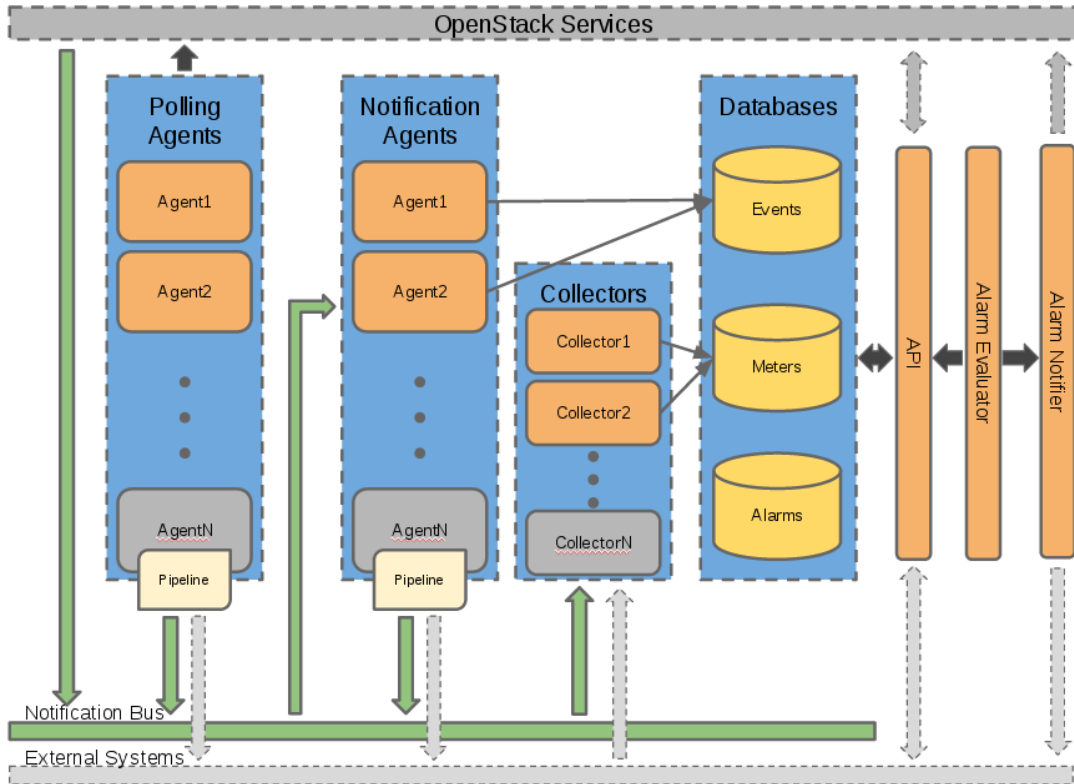
● ceilometer-collector writes to opentsdb socket via dispatcher



Gnocchi performance



Quick refresher on existing infrastructure



Ceilometer + Gnocchi integration

- Future plans:
 - transition Gnocchi API to Ceilometer API v3
 - move TSDB interaction to Ceilometer collector
- Steps we're doing right now:
 - integrate Gnocchi as separated piece of code inside Ceilometer workflow via Ceilometer Database Dispatcher mechanism

Covering existing ceilometer use-cases

- Alarming to drive Heat autoscaling, based on aggregating samples across all instances with matching metadata
 - *use cross-entity aggregation based on strongly-typed resource attributes, as opposed to free-from metadata*
- Reconstructing the resource state timeline, from per-sample resource metadata
 - *use queries over relatively infrequent events capturing state transitions*

No, we're not re-inventing TSDB here


- Existing specialized metrics-oriented DBs can be leveraged by Gnocchi's pluggable driver model
 - ***actively working on drivers for InfluxDB and OpenTSDB***
- Gnocchi itself provides a canonical storage driver based on Pandas and Swift
- In the specialized TSBD use-case, Gnocchi manages the resource-entity association & abstract archive policy concepts

Forward-looking questions to be resolved

- Consolidation of Gnocchi and Ceilometer codebases and core teams
- Migration of pre-existing datastores built up using classic Ceilometer
- Duration of deprecation path for the Ceilometer v2 API

Useful links

- <https://wiki.openstack.org/wiki/Gnocchi>
- http://bit.ly/gnocchi_blog
- http://bit.ly/gnocchi_docs
- http://bit.ly/gnocchi_changes



openstack™
CLOUD SOFTWARE



openstack™
CLOUD SOFTWARE